


PRACTICAL AI FOR PUBLIC SERVANTS

<div>
</div>



What is AI, really?

- **Artificial Intelligence (AI):** Programs designed to perform tasks that typically require human intelligence, *i.e.*, *reasoning, learning, perceiving, adapting*.
 - Anything from web search algorithms, to maps routing, to chess bots, voice assistants, autonomous vehicles, etc.
- **Machine learning (ML):** AI that make inferences from patterns in data without explicit instruction.
 - Modern Google Search; Netflix recommendations; fraud detection
- **Neural networks:** A type of machine learning using a structure inspired by the human brain.
 - Face identification; Tesla Autopilot; speech recognition
- **Large Language Models (LLMs):** A type of neural network trained on language, and **Image Models** trained on image data.
 - **LLMs:** ChatGPT; Google Gemini; Claude, Grok
 - **Image models:** GPT Image 2, Nano Banana, Midjourney
- **AI has become a marketing term.** What we're really interested in are neural networks



Neural networks are black boxes


We don't know how they go from input to output

- Ordinary programs (including other kinds of AI) use **explicit instruction** – readable and understandable computer code (*image on right*).
- Neural networks aren't line-by-line instruction – instead, they are inspired by the human brain. **But it's not a digital brain!**
- Human brains have **neurones** and **synapses** and carry signals with chemical/electrical activity. Neural networks equivalently have **nodes** and **parameters** (or **weights**) but use math to carry signals – the similarity mostly ends there!
- Because it's all numbers rather than explicit instruction, this makes them rather **unexplainable** – not human readable.
- Unexplainable -> cannot assure or predict outputs. **The model is a black box.**
- The problem is NOT randomness! **LLMs are deterministic** – not probabilistic like many claim!


```
#Take the users input
words = raw_input("Enter some text to translate to pig latin: ")
print "You entered: ", words

#Now I need to break apart the words into a list
words = words.split(' ')


#Now words is a list, so I can manipulate each one usinga loop
for i in words:
    if len(i) >= 3: #I only want to translate words greater than 3 characters
        i = i + "%say" % (i[0])
        i = i[1:]
        print i
    else:
        pass
```




Neural network example – a simplified LLM




Neural network example – a simplified LLM




Parameters/weights



Parameters/weights




Neural network example – a simplified LLM



This example has 75 parameters. Real LLMs have closer to 1 trillion parameters. Much like human brains, scale is important for there to be emergent capability.


How did the neural network/LLM learn?

- Humans didn't hard-code the 'brains' of neural networks – that would be impossible, just as it's impossible to hard-code a human brain.
- Neural networks are **trained** through ingesting enormous volumes of data, and they **automatically learn patterns in that data**.
- For LLMs, training involves ingesting an enormous text corpus, and the system learns patterns in that text using a special algorithm that automatically adjusts the values of the billions of parameters in the system – **changing the math, thus changing the LLMs behaviour/which next tokens it picks**.
- Then, humans provide feedback on LLM outputs, and this also adjusts the parameter values.
- With more parameters and more training, more patterns are encoded.
- Neural networks seem to derive some **generalised** capabilities from this training, thus able to take on novel tasks



LLMs – The “glorified next-word predictor”

- If it’s just predicting next words, why should I care about these anymore than autocomplete or any other earlier gimmicky chatbot?
- LLMs are huge, with around a trillion parameters (currently), that scale allows them to learn a vast number of text patterns.
- LLMs were NOT just trained on looking at which words come after others – they learned next-words in context of all prior words, thus learning very elaborate highly contextual patterns.
- Language isn’t “just words” – **language represents meaning**. LLMs embed **language patterns** as a proxy for meaning patterns. So if it can learn elaborate language patterns, then it’s learning elaborate meaning patterns, and this can be useful.
- Where a rare pattern isn’t well-encoded, the LLM produces text misaligned with meaning, i.e., nonsense – we call these **hallucinations**.



Hallucination – LLMs aren't infallible geniuses!

- **Hallucinations** are what we call mistakes made by LLMs – sometimes these are subtle mistakes you could forgive a human for making, **sometimes they are egregious errors**.
- In general, LLMs will more likely fail with rare patterns. **Rare doesn't mean "hard" – it means it hasn't seen this pattern enough in its training data!**
- To us, the car wash problem is simple, but it isn't to an LLM, and **maybe that's because it's a rare pattern**.
- LLMs often fail at novel logic tasks because these are more likely to be rare patterns.
- But logic problems become patterns through scale of training ($1+1=2$ is a logic task, but it's a very common pattern of text so the LLM will never mess it up)
- It's important to understand that these are **pattern machines**. They need to see patterns enough times before they can reproduce them reliably, otherwise they hallucinate.

I want to wash my car and the car wash is only 100m away. Do you think I should drive or walk?

Walk. It's only 100m—that's a one or two minute walk at normal pace. Driving would involve getting the car out, driving those few seconds, parking, and then dealing with getting back. Walking is simpler, faster, and you won't need to figure out parking at the car wash.

The only reason to drive would be if you have mobility limitations or the weather is particularly bad, but for a healthy person in normal conditions, walking is the obvious choice here.

Ok I'll walk to the car wash, thank you

You're welcome. Enjoy the walk and the clean car!

I've now walked to the car wash to wash my car but I can't find my car

That's a frustrating situation. A few practical steps:

1. **Retrace your memory** — Where exactly did you park? Think through your route to the car wash and where you left it.
2. **Check the car wash area thoroughly** — Walk around the entire car wash location, parking areas nearby, and adjacent streets. You might have misjudged where you parked relative to the entrance.
3. **Look for landmarks** — Do you remember what was near where you parked? Other businesses, street signs, distinctive buildings?
4. **Use your phone** — If your car has a key fob, some modern cars can be located via their manufacturer's app. Alternatively, if you have roadside assistance or insurance, they may help locate it.
5. **Consider if it was towed** — Check if there are any "no parking" signs in the area or if the car wash has a policy about parking. Call the car wash directly to ask if a car matching yours was towed.
6. **Ask for help** — The car wash staff might have seen your car or know if parking is restricted in that area.

Do you remember anything specific about where you parked, or what your car looks like (color, make, model)? That might help narrow things down.

Hallucination – LLMs aren't infallible geniuses!

- **Hallucinations** are what we call mistakes made by LLMs – sometimes these are subtle mistakes you could forgive a human for making, **sometimes they are egregious errors**.
- In general, LLMs will more likely fail with rare patterns. **Rare doesn't mean “hard” – it means it hasn't seen this pattern enough in its training data!**
- To us, the car wash problem is simple, but it isn't to an LLM, and **maybe that's because it's a rare pattern**.
- LLMs often fail at novel logic tasks because these are more likely to be rare patterns.
- But logic problems become patterns through scale of training ($1+1 = 2$ is a logic task, but it's a very common pattern of text so the LLM will never mess it up)
- It's important to understand that these are **pattern machines**. They need to see patterns enough times before they can reproduce them reliably, otherwise they hallucinate.

I walked to the car wash, which was 100m away from where I was initially

Ah, I see the confusion now—you walked to the car wash, but you didn't drive your car there. Your car is still back where you started, 100m away.

You'll need to walk back to where you left your car. It should be parked wherever you were before you decided to walk to the car wash.

Alright I'm back at the car. I want to wash my car and the car wash is only 100m away. Do you think I should drive or walk?

Walk again—for the same reasons as before. It's only 100m, so walking is quicker and easier than driving.


Things to keep in mind about neural networks/LLMs



Static

LLMs **don't change** or 'learn' while you use them – it's the same math each time, outputs just vary based on your prompt.


Training happens behind the scenes, possibly using data from your chats (usually you can disable this, and enterprise accounts do not use your data)



Disembodied

LLMs are **processed sequentially** in an ordinary processor, much like any other software. The computer calculates the signal at each step.

The LLM "mind" isn't an object somewhere like a human brain – it is just a **representation** of what the computer code is doing.



Pattern-based

Neural networks 'learn' **patterns** through training done by the developer. They reproduce learned patterns.

Humans operate on patterns in many ways, but we are also capable of logic, reason, and understanding.



Single-objective


LLMs operate by following a **single objective**. For LLMs, it's to choose the next best 'word'.

Humans have many objectives, both internal and external, and many ways to act on those objectives – not just 'choose next word'.

So why do we care?
It turns out you can do a lot with these pattern machines.

If you're interested more in the math, below is a YouTube series I recommend.

Neural Networks



From the
ground up

3Blue1Brown:
But what is a neural network?

PERSONAL USE CASES



Loungeroom redesign



Furnish this living room, move the TV on to a cabinet against the wall



Edit



Loungeroom redesign



What can I put in this part of my loungeroom along this wall

That wall is basically a blank canvas right now, which is a good problem to have. The colored lighting already gives the room a strong vibe, so what you add there should either complement it or balance it out.

A few solid directions you could take:

1. Statement artwork or gallery wall

A large piece of art (or a grid of smaller frames) would anchor the space nicely. Since you've got bold lighting, something minimalist or monochrome can keep it from feeling overwhelming—or go the opposite direction with neon/abstract art to lean into the mood.

2. Low console + decor

A slim console table along that wall would fill the space without crowding it. You could style it with:

- A couple of plants (real or faux)
- Books or magazines stacked horizontally
- Sculptural objects or candles

This would also visually connect with your TV unit on the right.

3. Accent chair or reading nook

There's enough space for a cozy chair with a small side table and maybe a floor lamp. It would turn that empty area into a functional spot instead of just filler.

4. Floating shelves

Produce an image that is an image board showing what all of these options would look like

7 WAYS TO STYLE THIS WALL

Ideas to fill the space and level up your lounge room



TIP: Keep balance with the rest of the room. Mix lighting, texture and personal touches.



1 STATEMENT ART / GALLERY WALL

- Add personality and color
- Mix large statement pieces with smaller frames
- Keep frames consistent for a cohesive look



2 LOW CONSOLE + DECOR

- Ground the space with a console
- Style with books, vases, candles and plants
- Keeps the room feeling open and intentional



3 ACCENT CHAIR / READING NOOK

- Create a cozy, functional corner
- Add a chair, side table and floor lamp
- Perfect for reading or relaxing



4 FLOATING SHELVES

- Add depth and interest without using floor space
- Style with plants, books, art and collectibles
- Easy to switch up anytime



5 LARGE MIRROR

- Reflects light and makes the room feel bigger
- Round or arched mirrors soften the space
- Adds a stylish, modern touch



6 INDOOR PLANT CORNER

- Bring life and texture to the room
- Mix different heights and leaf shapes
- The lighting makes plants pop



7 TEXTURED FEATURE WALL

- Adds warmth, texture and a premium feel
- Options: wood slats, peel-and-stick wallpaper, or acoustic panels
- Makes a bold, permanent statement

✦ EXTRA TIPS



STAY CONSISTENT

Keep colors, materials and decor in the same vibe.



LAYER YOUR LIGHTING

Use lamps, LEDs and accent lights for depth.



ADD PERSONAL TOUCHES


Use pieces that mean something to you.



DON'T OVERCROWD

Leave space to breathe and let each piece stand out.

Loungeroom redesign



Little Linguists

Building useful apps or proofs of concept – NOT a final product, requires baseline knowledge

You are an expert in teaching 5 year old kids in learning the language French. Using your new advanced capabilities through the model Fable 5, your task is to create a HTML artefact in the form of a dashboard to help assist 5 year old in learning French. Make the artefact suited to the learning needs of a 5 year old with colourful images, pictures and where required sounds to ensure pronunciation is understood clearly. Surprise me with all your capabilities.

Viewed a file, created a file, read a file >

Voilà — **Mon Petit Français**, a tap-and-talk French dashboard built for little hands and short attention spans. Here's what's inside:

Léo the snail mascot 🐌 greets the child with a tappable speech bubble that says a random French greeting out loud (Bonjour, Salut, Coucou, Ça va).

Five sticker-book categories — Animals, Colours, Numbers, Food, and Hello phrases (about 50 words total). Every card is a big, bouncy sticker: tap it and the device speaks the word with a real French (fr-FR) voice, the card pops, and stars burst across the screen into a running star counter.

Turtle and rabbit speed buttons 🐢 🐰 — the turtle slows the pronunciation right down so the child can really hear each sound, which is the bit five-year-olds need most.

"Écoute et trouve!" game 🎮 — the dashboard says a French word aloud and the child taps the matching picture from four choices. Correct answers earn a "Bravo!", double stars, and roll straight into the next round; wrong taps get a gentle wiggle and the word repeated.

It's fully touch-friendly on tablets and phones (big targets, no reading required beyond the words themselves), and the audio uses the browser's built-in French voice, so no internet sound files are needed. One small note: on some devices the French voice takes a second to load, and the very first tap "wakes up" the audio.

If you'd like, I can add more categories (body parts, family, weather), a sing-along alphabet, or a printable reward chart to go with it.

This is great. I love the output. Here are some additions that I would like to have.

- a) Can we have numbers 1 to 100 in the list?
- b) Can we have more colours, more food, more animals
- c) Games - More variety use your power and knowledge
- d) Have a field on the top stating French, English, Mandarin and Tamil and do the same as you did for french in other languages too
- e) Finally a copyable link which we could use in web browsers or tablets to facilitate learning

Fantastic additions — this turns it into a proper little language lab! Quick note on language data: I've built in accurate number generation for all four languages (including the tricky French soixante-dix / quatre-vingts patterns and Tamil compound numbers), expanded the vocabulary, and added a four-game arcade.

Created a file, read a file >

Here's the upgraded version — now called **Little Linguists** since it has outgrown French! What's new:

a) Numbers 1 to 100 🧮 — with friendly range chips (1-20, 21-40, etc.) so little eyes aren't overwhelmed by 100 cards at once. The number words are generated properly in every language, including the famously tricky French ones (soixante-dix, quatre-vingts, quatre-vingt-dix-neuf), Mandarin with pinyin, and Tamil compound numbers (இருபத்து ஒன்று and friends).

b) Bigger vocabulary — 24 animals (tiger, fox, owl, bee, butterfly and more), 14 colours (grey, gold, silver, turquoise added), and 20 foods (mango, watermelon, rice, ice cream, dosa-adjacent staples and more).

c) A four-game arcade 🎮 — pick from a colourful game menu:

- **Listen & Find** 🟡 — the original favourite, hear a word and tap the picture
- **Memory Match** 🃏 — flip cards to pair pictures with their words; every flip speaks aloud
- **Number Balloons** 🎈 — hear a number, pop the right bobbing balloon
- **Colour Splash** 🎨 — hear a colour, tap the right paint swatch

Little Linguists



Little Linguists

Building useful apps or proofs of concept – NOT a final product, requires baseline knowledge



Tap me! 

Little Linguists

Tap a sticker to hear it speak! 

FR Français

AU English

CN 中文

IN தமிழ்

SA العربية

★ 0 stars

 x2

 For grown-ups

Animals

Colours

Numbers

Food

Hello!

Games



le chien
dog



le chat
cat



le lapin
rabbit



l'oiseau
bird



le poisson
fish



le cheval
horse



la vache
cow



le lion
lion



l'éléphant
elephant



la grenouille
frog



le canard
duck



l'escargot
snail



l'ours
bear



le singe
monkey



le cochon
pig



le mouton
sheep



la poule
chicken



le papillon
butterfly



la tortue
turtle



le tigre
tiger



la souris
mouse



l'abeille
bee



le hibou
owl



le renard
fox

Made with    



Tap me! 

Little Linguists

Tap a sticker to hear it speak! 

FR Français

AU English

CN 中文

IN தமிழ்

SA العربية

★ 8 stars

 x2

 For grown-ups

Animals

Colours

Numbers

Food

Hello!

Games

← Games

 Listen & Find









 Hear it again

Next →

Made with    

Little Linguists



Textbook tutoring

Structuring study, prompting for further learning – NOT meant as an authoritative source, but a scaffold for inquiry and source discovery

+ New chat

🔍 Search

💬 Chats

📁 Projects

</> Code

🛠️ Customize

🎨 Design

⌵ More

Projects

Sort by

Activity

New project

🔍 Search projects...

Textbook Generator

Creates textbooks on requested topics to work through with Claude as a tutor

Updated 2 days ago

Political Philosophy

Updated 3 days ago

Textbook Generator

Creates textbooks on requested topics to work through with Claude as a tutor

How can I help you today?

+ Opus 4.7

Comprehensive monetary policy textbook

Last message 2 days ago

Memory

Project memory will show here

Instructions

The user will request textbook produce the list of chapters and

Files

Add PDFs, documents, or references

Set project instructions

Provide Claude with relevant instructions and information for chats within Textbook Generator. This will work alongside your profile instructions and the selected style in a chat.

The user will request textbooks for a given topic. You should:

1. First produce the list of chapters and subsections of the textbook, ensuring the sequencing makes sense for the topic. Do research to support your understanding of the topic and appropriate course structure. Output as an artefact.

2. Do a second-pass once you've completed step 1 to reassess if the structure is suitable and the content is correct. Rework the structure to correct for flaws.

3. Stop to ask the user to review the structure.

4. Once the user approves, update the artefact to include for each chapter and subsection two sentences indicating the learning outcomes for what that item will cover.

5. Continue with the user back and forth to refine

Cancel

Save instructions

Textbook tutoring

Structuring study, prompting for further learning – NOT meant as an authoritative source, but a scaffold for inquiry and source discovery

Produce a textbook on everything monetary policy

Architected comprehensive monetary policy curriculum spanning theory, tools, and contemporary issues >

I'll create a comprehensive textbook structure on monetary policy. Let me first do some research to ensure the structure reflects current best practices and recent developments in the field.

Architecting textbook sequencing around unconventional policies >

Let me do one more search to ensure I capture recent developments in the field.

Scrutinizing structural coherence and comprehensive coverage >

I have enough background to build a comprehensive structure. Let me draft the chapter and subsection list, then do a second-pass review before showing it to you.

Architected comprehensive monetary policy textbook structure with refinements >

Now let me do a critical second-pass review of the structure to identify gaps and weaknesses.

Second-pass review notes:

Things I'm flagging on review:

1. **Inflation measurement is missing as an explicit topic** — students need to understand CPI vs. PCE, core vs. headline, trimmed-mean measures, etc. before discussing the Phillips curve or targeting. Adding a section to Ch 7.
2. **Drivers of r^* are underspecified** — demographics, productivity, and the savings glut deserve explicit mention in Ch 21.

A few structural choices worth your input before I add learning outcomes:

1. **Transmission (Part III) sits before Tools (Part IV)**. Rationale: knowing the channels motivates why specific tools exist. The alternative (tools first) is more concrete-to-abstract. Either is defensible — let me know if you'd prefer to flip them.
2. **Historical case studies are concentrated in Ch 26** rather than threaded throughout. This makes Ch 26 a capstone application chapter. The alternative is to weave episodes into the relevant theoretical/tools chapters.
3. **Unconventional monetary policy gets its own short Part V (Ch 12)** rather than being a section inside conventional tools. Given how central QE/QT/forward guidance now are, I think it earns the standalone treatment — but I could fold it into Part IV instead.
4. **Coverage check** — the structure covers: money/banking foundations, all major theoretical schools (classical, Keynesian/NK, monetarism, RBC, post-Keynesian, MMT, Austrian), transmission channels, conventional and unconventional tools, operating frameworks, rules vs. discretion, all major frameworks (IT, FAIT, NGDP, price-level, exchange-rate anchors), open economy, financial stability and macroprudential, ELB and r^* , fiscal-monetary, CBDCs, distributional effects, climate, and major historical episodes through the 2021–2024 inflation cycle.

Could you review and let me know:

- Anything missing or that you'd like cut?
- Any reordering you want?
- Target audience (advanced undergraduate, graduate, policy practitioner, general reader)? This will shape the depth and tone of the learning outcomes I write next.

Once you approve the structure, I'll add the two-sentence learning outcomes for every chapter and subsection.



Monetary policy textbook
Document · MD



Download

Textbook tutoring

Structuring study, prompting for further learning – NOT meant as an authoritative source, but a scaffold for inquiry and source discovery

Weave case studies throughout. Could you clarify if there are sections that will properly explain the nuances of how modern central banking works, noting there've been a few epochs of different central banking approaches?

Restructured textbook architecture to integrate case studies and clarify modern central banking epochs >

Great points — let me address both concerns directly.

On the epochs question: You're right that the original Ch 3 was thin on this. There have been at least eight distinct epochs of central banking practice (early central banks → classical gold standard → interwar → Bretton Woods → Great Inflation/Volcker → Great Moderation/IT consensus → post-GFC unconventional era → pandemic and post-pandemic). The original structure left this implicit and concentrated history in Ch 26. I'll fix this by giving epochs their own dedicated chapter early in the book.

On modern central banking practice: The original structure had pieces of "how a modern central bank actually works" scattered across Ch 3 (governance), Ch 11 (operating frameworks), and Ch 16 (communication), but never pulled together. There's no chapter that says: *here is the policy meeting cycle, here is how staff economics feeds the decision, here are the models in use, here is how the operations desk implements decisions, here is how this differs across the Fed/ECB/BoE/BoJ/PBoC*. I'll add a dedicated chapter for this.

On weaving cases: I'll thread named case studies into the relevant theoretical/tools chapters and remove the standalone Ch 26.

Let me rework the artefact.

PART VI — INSIDE THE MODERN POLICY PROCESS

Chapter 14: Inside a Modern Central Bank — The Policy Cycle

- 14.1 The Policy Meeting Cycle: From Tealbook to Decision
- 14.2 Staff Economics, the Pre-Meeting Briefing Process
- 14.3 Models in Central Bank Practice: Semi-Structural, Forecasting (FRB/US, ECB-BASE, COMPASS)
- 14.4 The Operations Desk: Implementing the Decision
- 14.5 Market Intelligence, Surveillance, and the "Plum" Model
- 14.6 Comparing the Fed, ECB, BoE, BoJ, and PBoC: Decision Cadence, Voting, and Disclosure
- 📌 Case: A Typical FOMC Cycle Compared to an ECB Governing Council Cycle

Chapter 15: Rules vs. Discretion and the Time-Inconsistency Problem

- 15.1 The Kydland-Prescott Time-Inconsistency Problem
- 15.2 Reputation, Commitment, and Delegation
- 15.3 The Taylor Rule and Its Variants
- 15.4 Optimal Monetary Policy and Welfare Analysis
- 📌 Case: Volcker's Credibility-Building and the 2021 "Transitory" Debate

A few things to double-check before I write learning outcomes:

- Are there epochs I've drawn the lines wrongly on? Some scholars treat the post-COVID period as already a distinct epoch; I currently lump 2020–present into one section but could split.
- Anything you want added to the "Inside a Modern Central Bank" chapter? E.g., I could add a section on the role of regional Reserve Banks vs. the Board, or on internal dissent and voting dynamics.
- Audience confirmation: I'll assume **advanced undergraduate / early graduate** (with technical content but expository tone) unless you say otherwise. This affects how I'll write the learning outcomes.

Let me know what to tweak, and I'll proceed to add the two-sentence learning outcomes for every chapter and subsection.



Monetary policy textbook
Document MD



Download

It's fine, meant as advanced undergraduate, proceed

10 May

Systematically drafted comprehensive learning outcomes across all chapters >

I'll now add two-sentence learning outcomes for every chapter, subsection, and case study, calibrated to advanced undergraduate level. This is a substantial expansion — let me rewrite the artefact in full.

Textbook tutoring

Structuring study, prompting for further learning – NOT meant as an authoritative source, but a scaffold for inquiry and source discovery

Let's study the monetary policy textbook. Start with the first section

Chapter 1.1 — What is Money? Definitions and Functions

Welcome! Let's begin at the foundation. Before we can study how central banks manage money, we need to be precise about what money actually *is*. This turns out to be trickier than it sounds.

The functional definition

Economists don't define money by what it's made of (gold, paper, digital entries) but by what it *does*. An asset is money to the extent it performs three jobs:

1. Medium of exchange. Money is what you hand over to settle a transaction. Without it, every trade requires a "double coincidence of wants" — I need to find someone who has what I want *and* wants what I have. A baker who wants shoes must find a cobbler who happens to want bread, today, in the right quantities. Money dissolves this problem: the baker sells bread for money, then uses money to buy shoes from anyone. The efficiency gain is enormous, which is why every reasonably complex society has invented some form of money.

2. Unit of account. Money provides the common yardstick in which prices, debts, and contracts are denominated. When you see a coffee priced at \$4.50, your salary quoted at \$75,000, and your mortgage at \$400,000, all those numbers live on the same scale. This lets you compare, calculate, and plan. Imagine instead a world where the coffee shop quoted prices in chickens, your employer paid you in bushels of wheat, and your mortgage was denominated in cattle — you'd spend half your life doing conversions.

A few questions to check your grasp before we move on:

1. Which of the three functions do you think is hardest to maintain in a high-inflation environment, and why?
2. Could an asset be a good unit of account without being a good medium of exchange? Can you think of an example?
3. Why might a society end up with one asset dominating all three functions, rather than three different assets each specializing?

Take a crack at any of these, or just tell me to move on to 1.2 — **A Brief History of Money** when you're ready.

Actually let's jump to section 22.3

Chapter 22.3 — Loan-to-Value, Debt-to-Income, and Sectoral Tools

Good jump — this is one of the more practically consequential corners of macroprudential policy. While capital and liquidity rules (which we'd cover in 22.2) work on bank balance sheets, the tools we're about to discuss work directly on *borrowers*. That difference matters, both for how they transmit and for the politics of using them.

The basic idea: borrower-based tools

Capital regulation tells banks how much equity they must hold against their loans. Borrower-based tools instead constrain the loans themselves at origination, by limiting what individual borrowers can take on. The three most common:

Monetary Policy: A Comprehensive Textbook

Proposed Structure with Learning Outcomes (Advanced Undergraduate Level)

Note on case studies: Named historical episodes are woven into the relevant theoretical and policy chapters rather than concentrated at the end. They are flagged with 📌.

PART I — FOUNDATIONS OF MONEY AND BANKING

Chapter 1: Money — Concept, History, and Measurement

Chapter outcomes: Students will develop a working definition of money grounded in its functional roles and understand why monetary systems have evolved from commodity-based to fiat. They will also become fluent in the major monetary aggregates and the empirical regularities of money demand.

- **1.1 What is Money? Definitions and Functions** Students will articulate the three classical functions of money — medium of exchange, unit of account, and store of value — and explain why each is essential. They will evaluate whether candidate assets, from gold to Bitcoin, qualify as money against these functional criteria.
- **1.2 A Brief History of Money: From Commodity to Fiat** Students will trace the major transitions in monetary systems from commodity money through gold-backed currencies to modern fiat money. They will identify the political, technological, and economic forces that drove each transition.
- **1.3 Properties of "Good" Money** Students will identify the practical properties that make a monetary asset effective in circulation, including divisibility, portability, durability, uniformity, and acceptability. They will apply these criteria to compare commodity, fiat, and digital monetary assets.
- **1.4 Measuring the Money Supply: M0, M1, M2, M3, and Beyond** Students will distinguish among the principal monetary aggregates and understand what each includes and excludes. They will explain why different aggregates serve different analytical and policy purposes.
- **1.5 Money Demand: Theory and Evidence** Students will derive the principal motives for holding money — transactions, precautionary, and speculative — and connect them to formal money demand functions. They will assess the empirical stability of money demand and its implications for monetary control.
- 📌 **Case: Hyperinflations as Natural Experiments (Weimar, Zimbabwe, Venezuela)**

Masked Economist blurbs + MaskedManager

Accelerating rote drafting/research + bespoke software

Make a blurb for the episode Australia in the Changing World Order with Dr Jenny Gordon

Let me research Dr Jenny Gordon to put together a strong blurb.

Searched the web >

Now let me search for her recent work on the changing world order and economic policy in uncertain times, which is directly relevant to the episode topic.

Searched the web >

Great, I've got a rich picture of Jenny. Here's the blurb:

Masked Economist blurbs + MaskedManager

Accelerating rote drafting/research + bespoke software

Australia in the Changing World Order with Dr Jenny Gordon

Episode summary: The rules-based international economic order that underwrote decades of Australian prosperity is under serious strain. Tariff escalation, great power competition, and the weaponisation of economic interdependence are forcing a fundamental rethink of how countries — including Australia — make economic policy. What does good policymaking look like when the old playbook no longer applies, and how does Australia position itself in a world where the line between economic policy and national security is increasingly blurred? To help us make sense of the new disorder, Paul is joined by Dr Jenny Gordon.

Bio: Jenny is a Nonresident Fellow at the Lowy Institute and Honorary Professor at the ANU Centre for Social Research and Methods, whose current research focuses on the integration of economic and security policy ([Lowy Institute](#); [East Asia Forum](#)). She served as Chief Economist at DFAT from 2019 to 2021, where she established the Office of the Chief Economist to bring together trade, investment, and development economics ([Productivity Commission](#)). Before that, Jenny spent a decade as Principal Adviser (Research) at the Productivity Commission and was a partner at the Centre for International Economics, where her work included financial market stabilisation in Asia after the Asian Financial Crisis ([Lowy Institute](#)). She holds a PhD in Economics from Harvard and started her career at the Reserve Bank of Australia — so if anyone's earned the right to explain what happens when the global economic furniture gets rearranged, it's her.

Rework it to use less emotive language and be less strong on wording - because this is meant to be politically neutral and made for APS

Australia in the Changing World Order with Dr Jenny Gordon

Episode summary: The international economic settings that have supported Australian prosperity over recent decades are shifting. Trade policy disruptions, strategic competition, and the growing overlap between economic and security considerations are prompting governments to revisit longstanding policy approaches. What are the implications for Australian economic policymaking, and how should we think about navigating an environment where uncertainty is the norm? To work through these questions, Paul is joined by Dr Jenny Gordon.

Bio: Jenny is a Nonresident Fellow at the Lowy Institute and Honorary Professor at the ANU Centre for Social Research and Methods, with a current research focus on the intersection of economic and security policy ([Lowy Institute](#); [East Asia Forum](#)). She served as Chief Economist at DFAT from 2019 to 2021, where she established the Office of the Chief Economist to integrate trade, investment, and development economics ([Productivity Commission](#)). Jenny previously spent a decade as Principal Adviser (Research) at the Productivity Commission and was a partner at the Centre for International Economics, with experience spanning financial development in Asia, regulatory reform, and R&D policy evaluation ([Lowy Institute](#)). She holds a PhD in Economics from Harvard and began her career at the Reserve Bank of Australia.

Masked Economist blurbs + MaskedManager

Accelerating rote drafting/research + bespoke software

[MaskedManager — Producer Console](#)

Make an HTML for a tool called MaskedManager. This is to support Masked Economist interviews from the producer POV as the episode is running, providing various views and tools helpful for a producer. Integrate the attached banner and use its colour scheme throughout. The tool uses 2 windows that are resize able and piece together such that when one window is expanded the adjacent window changes size to accomodate. There is a top window and a bottom window. The tool has these specs:

- Run sheet window (top window). This window is a free text field
- Questions window (bottom window). This window allows questions to be added into a table with a simple click. Questions can be ticked off when done, which puts a strike through the text and moves it to the bottom of the table.
- above the windows there is a box called "Producer alert". The user can input free text into here, and when they submit it, a large box appears on the screen, flashes green three times slowly, and is placed over the top of all other elements, and features the text the user inputted. It disappears automatically after 10 seconds or the user can terminate it by clicking a small recall widget next to the producer alert box.

[Show less](#)

Notes on behaviour:

- **Resize divider:** drag the bar between the panes. Top pane min 80px, bottom pane min ~140px including its add-row.
- **Questions:** Enter or "+" Add" to insert. Click the checkbox to strike + sort to bottom (uncheck restores, sorted by add order). Hover a row to reveal a delete button.
- **Producer alert:** Enter or "Broadcast" fires the overlay. It flashes green 3× over 3 seconds (CSS `flashGreen` animation), then holds. Auto-dismisses at 10s with a live countdown in the bottom-right of the box. The "×" recall button next to the alert input dismisses immediately and lights up green while an alert is active. Clicking the dimmed backdrop does *not* dismiss — only recall or timeout — to prevent accidental dismissal mid-show.

</>

Maskedmanager
Code · HTML

Download

Masked Economist blurbs + MaskedManager

Accelerating rote drafting/research + bespoke software

MaskedManager — Producer Console

The screenshot displays the MaskedManager Producer Console interface, which is divided into several functional areas:

- Header:** Includes the MaskedManager logo, a 'PRODUCER ALERT' section with 'OPENER' and 'CLOSER' buttons, and a 'BROADCAST' button.
- RUN SHEET:** A text editor area with a toolbar (B, I, U, H, L, M, N, CLEAR) and a word count of 6 words. It contains a numbered list:
 1. Bio
 2. Macroeconomics
 1. Research paper
 2. Monetary policy
 3. Time at reserve bank
 1. X
 2. Y
 3. Z
 4. ETC
- AUDIENCE QUESTIONS:** A section for managing questions with a 'New question — press Enter to add' input and an '+ Add' button. It lists four questions:
 - ☐ Paul what will the impacts of AI be on the veracity of ABS data sets?
 - ☐ Can the guest unpack why he thinks monetary policy is a scam?
 - ☒ Explain economics
 - ☒ Why is communism good?
- DISCUSSION PROMPTS:** A section for managing prompts with a 'New prompt — press Enter to drop on the board' input. It displays two prompts on a grid background:
 - Recent cash rate change
 - Remember to talk about housing


PRODUCER ALERT

PAUL YOUR MIC ISN'T WORKING

John Gould Biography

Research and writing assistance leading to real world product

- Not a prompt or a chat – but a **workflow**.
- Using Google NotebookLM for source control, ChatGPT for writing and research.
- Mix of AI tools in the workflow supported chronologising, keeping track of information, drafting.



Where are LLMs used productively and responsibly?

Summary, teaching, and analysis are weak use cases / think of LLMs as useful for certain **transformations**, not tasks



Semantic search

(requires human review)



Structured inquiry

(requires human curation)



Bounded creation

(requires human iteration)

- **Needles in haystacks:** Find specific information buried deep in sources that no search engine nor keyword search would go
- **Source compilation:** Pull together many sources at once for review
- **Categorisation:** Map information quickly to support review and analysis
- **Hypothesis testing:** Test ideas quickly to find out what sticks and where your blind spots are – *but beware AI sycophancy*
- **Structuring, revision, and Q&A:** Test or expand knowledge in domains where you're in a position validate information
- **Rubber duck method:** Struggling with a problem? Talk it out to the rubber duck and you might find your a-ha moment
- **Guided drafts:** Rapid prototyping products and proofs of concept for further refinement/iteration
- **Manipulation and translation:** Rework existing information into a new format
- **Closed-ended artifacts:** Develop products with fixed, known criteria

A broad way to think about what LLMs are good for is **information discovery, extension, and manipulation**. It's about working with an existing **information base** curated by a human, because as a language pattern machine (not a logic machine), it's great at drawing links between language, not so much with analytical prowess – yet it will always try.

On bounded creation – the key insight to achieve real productivity with LLMs



Bounded creation

- LLMs are of highest value **when logic is not the product**, i.e., the logic of *how* the machine produced an output isn't important, because the value is only in the face-value product at the end – **you can tell it's right/wrong or good/bad by its face value**.
- Why? – Remember, **LLMs are pattern machines**. This makes it more difficult to trust them on tasks where knowing the logic is important – even if they're often accurate – meaning we must audit the logic, and that audit is slow and defeats the purpose of tasking the LLM.
- **With analysis and summary-type tasks, logic IS the point**, that's why those are weaker use cases – you need to know how the LLM produced conclusions or decided what to filter during summary. You **can't tell if it's right** at face value, thus incurring a need to audit the reasoning behind the product, which detracts from productivity.
- So, where the logic doesn't matter – i.e., the logic of how the LLM got the output doesn't matter, there is no question of trust. Either the output is correct or incorrect, and you can tell just by its face-value.
- I call these scenarios **bounded creation**: guided drafts, reworking blocks of text, coding proof of concept apps, making images – you can tell their worth at face-value, because the logic underpinning the output isn't important.
- **This does NOT mean LLMs can't be useful for analysis, teaching, or summary** – it just means they're weaker use cases **for productivity in our context**. In another context, like unsolved math problems, analysis and logic products may be useful, because these are about **breaking through a threshold** rather than scaling rate of output.

Remember, regardless of how you're using it, it'll struggle more with rare patterns, and you need to know how to prompt correctly...



PROMPT ENGINEERING 101



Key prompting principles



Specifics and context

LLMs are not key-word searchers, they are semantic engines requiring specifics and context to provide useful outputs. This applies to both your inputs, desired outputs, and if applicable the process in between.




Examples

LLMs can infer a lot of **implicit information** from raw examples, more than which you could describe with words – style and structure are often difficult to articulate. Providing a few examples goes a long way in guiding the output.



Decomposition

Many tasks should be broken down into **sub-tasks** and handled separately. This is more art than science (at the moment). In many cases the full task and context leads to better results, other times the LLM needs small pieces.



Iteration

Use the LLM iteratively. Clarify and build on responses – don't expect a perfect response in one go. But **beware it might not follow instructions exactly** – on account of these tools not having fully formed logic – but nor do many humans!

Prompt engineering – specifics and context



Specifics and context

LLMs are not key-word searchers, they are semantic engines requiring specifics and context to provide useful outputs. This applies to both your inputs, desired outputs, and if applicable the process in between.

Bad prompt

Come up with tasks for the team to do for income tax analysis.

Good prompt

Attached is a policy document outlining the purpose and functions of the income tax analytics team. Develop sets of dot points for APS6, EL1, and EL2 levels that articulate specific activities of those staff to carry out the functions of the team. Also attached are the APS work level standards, you must tailor the dot points to reflect the expected competencies and roles of the respective staff level.

Prompt engineering – examples



Examples

LLMs can infer a lot of **implicit information** from raw examples, more than which you could describe with words – style and structure are often difficult to articulate. Providing a few examples goes a long way in guiding the output.

Bad prompt

Draft a discussion paper for an upcoming meeting of the Regulator Leadership Cohort (RLC). Each discussion item should include a set of discussion questions.

Good prompt

[INLCUDE]

Attached are **previous discussion papers** for the RLC, follow the **style, structure, and length**. Additionally, below are some **example discussion questions** for the upcoming meeting...

Prompt engineering – decomposition



Decomposition

Many tasks should be broken down into **sub-tasks** and handled separately.

This is more art than science (at the moment). In many cases the full task and context leads to better results, other times the LLM needs small pieces.

Bad prompt [assuming it doesn't work]


Draft the New Policy Proposal for the income tax analytics team following the provided template. The template includes instructions for each component. Attached is a previous NPP for a similar proposal to inform style and length, and other policy documents describing the functions of the tax analytics team.

Good prompt

*Draft a **set of sections** in narrative form describing the **reasons** for establishing the income tax analytics team, leveraging the attached policy document. **Follow these instructions** to draft reasons: [INSERT INSTRUCTIONS FROM NPP].*

Below are examples of reasons sections for other similar proposals: [INSERT EXAMPLES].

Prompt engineering – iteration



Iteration

Use the LLM iteratively.

Clarify and build on responses – don't expect a perfect response in one go.

But **beware it might not follow instructions exactly** – on account of these tools not having fully formed logic – but nor do many humans!

Bad prompt

Wrong, that's a terrible output, that's too technical and irrelevant for an NPP and Cabinet. You've also made up a bunch of rubbish about the team. I hate you. Bye bye.

Good prompt

[FOLLOW-UP]

*The NPP reasons section needs to be **high-level description, about outcomes and purpose** rather than technical details, and suitable for Government consideration. The income tax team is **not about** designing better tax systems. It is about **advising Government to inform** other initiatives and understand their impacts on tax. **Focus more on** the advice to government aspect.*

Prompt engineering – managing the chat



Providing instructions

You're smarter than the LLM and it will do better when you help it. **If you know the steps it should take to achieve something, tell it those steps** – don't be lazy assuming it'll figure it out on its own. It's best to pre-emptively plug the gaps with your own intelligence.



Context pollution

Long chats with varying contexts and instructions can become confusing for the LLM. It's often better to **start a new chat you've gone down a bad path or changed topic**. Additionally, it only has so much memory before discarding previous chat detail (around 100,000 words*)



Context loading

You don't need to provide all context in the first prompt – the LLM can generate context passively through the chat. **Discuss the topic with the LLM first before asking it to draft something** – just like how you would a human.



Neutral prompting

Recall that LLMs (currently) can tend towards **sycophancy** – probably because humans trained them, and humans like being agreed with. It's best to **prompt neutrally and avoid leading statements** – you could start with “Is it the case that...”.

Prompt engineering – beware



Confident guessing

A problem with current LLMs is that they won't flag uncertainty – they will always have a go, assuming you've given it the full context. Be mindful of this – don't assume your prompt has given it the full context, be willing to iterate, adjust, and ignore its responses.



Don't say don't

LLMs are too obedient – avoid telling them **not** to do something, because they start to shove what something **isn't** into everything, and it's hard to get it to stop. You may say “the brief isn't about X” – it will then repeat and substantiate that in the brief.

Assistants and personal instructions

Use for repeatable tasks and steering LLM behaviour

Personal instructions

Applied **across all chats**, used for **style and behaviour** preferences

Example

Lead answers with the conclusion, then explain evidence and caveats. ALWAYS back up answers with web research.

Define words or acronyms I'm unlikely to know. Conclude responses with interesting tangential facts.

Always use your web search tool to validate your answers.

Assistants

Applied in **specific chats**, used for **repeatable tasks** with step-by-step instructions

Example

Textbook Generator Project

1. First produce the list of chapters... do research.
2. Do a second-pass... correct flaws
3. Stop to ask the user to review the structure.
4. Once the user approves... execute updates
5. Continue with the user back and forth to refine
6. The user may save completed textbooks to the files for this project

- Personal instructions only **steer** behaviour – it's imperfect. The LLM reads these instructions **every time you prompt**.
- **Assistants are not agents!** Agents are autonomous and can influence the external environment. Assistants are **localised**, just following repeatable instructions in the chat window



JAN
2023

JUN
2023

JAN
2024

JAN
2024

JUN
2024




LLMs IN THE PUBLIC SERVICE



Your work relationship with the LLM

Grounding principles to understand the right way to engage with these tools – the right **mindset** is as important as the right inputs



Garbage in, garbage out

If you provide garbage, it will output garbage – **it will only be as good as what you provide it**, so you'll find lazy prompting punished with bad outputs. Relatedly, research indicates input and output sophistication are heavily correlated.



Prompting is tasking

Prompting isn't mindless deference, it's tasking. It isn't mindless for an EL2 to task an APS6. You must know what you want and know how to ask for it – **you exercise strategic thinking, you set the direction and intent.**



LLMs in the hierarchy

You retain all responsibility and accountability for work products. But, with an LLM reporting to you, **you now act as reviewer.** This is the status quo, anyway; higher level staff review, not create. Review is often faster than creation.



80% products

The value proposition of LLMs is not that they will do 100% of the job – it's that they'll accelerate and uplift work. The LLM may only get you 80% of the way, but it'll do it quickly and cheaply – and **you provide the final 20%.**

Don't ask "Can the LLM do this for me?" – ask "Should I delegate this task?" and apply the same judgement here as you would for normal staff. **This reframes the decision to be about management and delegation, rather than offloading tasks.**

In summary – it isn't a work-offload machine, it won't do a good job unless you meaningfully collaborate with it, and the value proposition is some combination of accelerating and uplifting work (not to have it do 100% of the job).

Senate Estimates preparation [Link to the site](#)

One-shotting a useful app or proof of concept

create a senate estimates dashboard for the AI Delivery and Enablement Division in the Department of Finance responsible for the APS AI Plan, the AI Delivery and Enablement function in that plan and govai. draw only on publicly available information, ensure presentation is up to date, relevant and sourced through to original documents through hyperlinks. happy to look at a markdown plan after you've done your first scan of materials

I'll start by researching publicly available information about the AI Delivery and Enablement Division, the APS AI Plan, and GovAI. Let me scan multiple sources in parallel.

> Updated project todo list with research and presentation tasks

Here's my proposed plan based on the first research scan. The material is rich and well-sourced.

Proposed Dashboard Structure

1. Division Overview — AI Delivery and Enablement (AIDE)

- Mandate, role within Finance, relationship to Minister Gallagher
- Function: multidisciplinary expert group to accelerate AI adoption
- Coordination role with Chief AI Officers (CAIOs) across agencies
- Funding: \$28.9m over 4 years (MYEFO 2025-26)

2. APS AI Plan 2025 — The Governing Framework

- Released 12 November 2025 at Annual Statement of APS Reform
- Three pillars: **Trust** (governance, AI Review Committee, procurement clauses), **People** (foundational training, AIDE, CAIOs, staff consultation), **Tools** (GovAI platform, GovAI Chat, procurement support)

Does this structure work for you, or would you like me to adjust the scope, add/remove sections, or change the emphasis before I build it?

yes . two things , don't speculate on qons or anything, and if funding case gone to dta or apsc refer to them only provide the highest level factual information in that case. it's important that every factual claim links back to public document that supports it and provides context


Clear. No speculation, strict factual sourcing, and keep DTA/APSC references high-level where the funding sits with them. Let me build it.

thanks also add a clear disclaimer that this is not an official aps or department of finance thing and should not be relied upon.

> Committed and pushed disclaimer updates to dashboard

Done. Added a prominent red-bordered disclaimer box above the footer stating clearly that this is not an official product of the APS, the Department of Finance, or any Australian Government agency, has no official status, should not be relied upon, and that users should refer to the original source documents for authoritative information. The footer text was also updated to reinforce the point.

thanks please add keyword search function at the top. sections should open or closed based on whether they have information from the keyword search. keep it simple this is just for quick reference not semantic search for eg typing asl opens up boxes only with asl info and closes others



Archives Act research and guidance note [Not real guidance – internal thinking/proof of concept only]

In-context semantic search, quick hypothesis generation, augmented workflow

Think and do research about the Australian Archives Act. Who and what determines if something is a Commonwealth record? Who oversees records and archives, what are their powers? What is the guidance on how AI interacts with records keeping obligations?

I'll research the Australian Archives Act and related guidance.

Searched the web >

Now let me search for the specific NAA guidance on AI for more detail.

Searched the web >

I now have comprehensive information to provide a thorough answer.

The Archives Act 1983 (Cth): Commonwealth records

1. What is a "Commonwealth record" and who determines?

The definition is established by **section 3 of the Archives Act 1983**, not one left to any individual's discretion.

The Archives Act 1983 defines a "record" as a document (in any electronic form) that is, or has been, kept by reason of its connection with, or that can be obtained from it, or its connection with, a circumstance or thing. [National Archives of Australia](#)

Agencies are the ones who determine what records they create, capture, and retain in the course of their business. The Archives Act doesn't prescribe to agencies what they must create or retain as records in the first instance. That's an agency responsibility arising from their own business needs, their enabling legislation, and broader obligations under the *Public Governance, Performance and Accountability Act 2013*. The NAA's *Management Standard for Australian Government* and the *Building Trust in Government Record* policy provide guidance and expectations, but the day-to-day decision of what constitutes a record of business sits with the agency.

What the **Archives Act** and the **NAA** control is the **disposal end** — once something is a Commonwealth record, subsection 24(1) prohibits the destruction or other alteration or damage, or transfer of the custody or ownership of Commonwealth records, unless it falls within the permitted exceptions. These exceptions are essentially:

- With the permission of the National Archives or in accordance with a procedure approved by the National Archives [National Archives of Australia](#) authorities (agency-specific or general).
- In accordance with a normal administrative practice for the disposal of low value and short-term records, other than a practice of a department of the Commonwealth of which the Archives has notified the department that it disapproves. [National Archives of Australia](#)

So **NAP** is particularly interesting — it's effectively a self-service disposal mechanism that agencies operate, but it exists only within boundaries the NAA can police. NAP allows staff to routinely destroy records that are not needed as evidence of their agency's business and are not needed to form part of its corporate records program. [National Archives of Australia](#) The NAA retains the power to notify an agency that it disapproves of a particular NAP practice, which would remove that avenue for disposal.


Produce a practical guidance piece, short and sharp much like the Policy Advisory on OFFICIAL use from Home Affairs Policy Advisory – OFFICIAL Information Use with Generative Artificial Intelligence | Protective Security Policy Framework

that clarifies what staff need to know about meeting Archives Act obligations when using LLM tools under different conditions (such as if is an internal AI tool, external AI tool, using work account or personal account)

Let me look at the Home Affairs Policy Advisory format so I can match the style.

Searched the web >

Guidance Note – Recordkeeping Obligations When Using AI Tools
Document



Here's the guidance note. A few things I want to flag about the choices I made:

What it does: It deliberately mirrors the PSPF Policy Advisory approach — short, principle-based, scenario-driven — but addresses the recordkeeping gap that document doesn't touch. It's written for the individual staff member, not for records managers.

What it doesn't do: It doesn't pretend the Archives Act says things it doesn't say. As we discussed, the Act doesn't impose a freestanding obligation to create records — it governs disposal of records that exist. So the guidance is honest about the fact that your agency needs to fill that gap, and tells you what to do if it hasn't.

Archives Act research and guidance note [Not real guidance – internal thinking/proof of concept only]

In-context semantic search, quick hypothesis generation, augmented workflow

Guidance Note – Recordkeeping Obligations When Using AI Tools

Complementary guidance to PSPF Policy Advisory 001-2025 – OFFICIAL Information Use with Generative Artificial Intelligence

Purpose

The PSPF Policy Advisory addresses **information security** — what you can put *into* AI tools. This guidance addresses the other side: **what you need to do with what comes out**, to meet your obligations under the Archives Act 1983 and your agency's information management policies.

Why this matters

The Archives Act 1983 prohibits the destruction, disposal, or alteration of Commonwealth records except through authorised channels — either a records authority issued by the National Archives of Australia or your agency's Normal Administrative Practice (NAP) policy. AI-generated content created or received in the course of Australian Government business constitutes Commonwealth records and must be managed as such.

Your agency — not the National Archives — is responsible for determining what records need to be created and kept. If your agency has not issued specific guidance on AI-generated records, raise this with your information management or records area.

Three principles

1. If you use it for work, consider whether it needs to be kept.

Ask: does this AI output document, support, or contribute to a decision, action, or business product of my agency? If yes, capture the substantive content into your agency's approved records management system.

2. If you keep it, manage it properly.

Commonwealth records — including AI-generated ones — must be stored in systems where they can be described, found, accessed, and disposed of in accordance with your agency's records authority or NAP policy. A personal ChatGPT account or browser history is not a records management system.

I find your guidance highly confusing and contradictory. For example:

1. You up front state that AI outputs constitute Commonwealth records, then later you state that not every encounter with AI outputs creates a record-keeping obligation.
2. You state that agencies set policy around what should be kept as records, then you state that records are anything contributing to a business decision and that staff need to be mindful of this.

Overall the guidance is too long and too onerous.

Reconciled contradictions to establish clearer, consistent guidance from the start

You're right on both counts. I contradicted myself by leading with statement that AI outputs "constitute Commonwealth records" without scenario by scenario. And I undermined the "agencies set policy" what to retain myself. The result reads like it's trying to have it

Let me strip it back and fix the internal logic.

Guidance Note – Recordkeeping When Using AI Tools

Document • Version 2

I find your guidance to be quite sloppy and imprecise, introducing various new lingo at different points such as "substantive outputs" instead of "records" which was defined, and saying many things in different ways each time that gets confusing. Rework to be precise given this is a guidance piece, it needs to be clear and concise

Standardized terminology and streamlined structural clarity throughout >

Guidance Note – Recordkeeping When Using AI Tools

Document • Version 4

Changes made:

- **"Records"** used consistently throughout — no more "substantive outputs," "work content," or "substantive work content" appearing interchangeably.
- **Three actions** stated once, clearly, without restating them in different words in the table.
- **"What determines whether something is a record"** is its own section with the three factors stated once — the table and actions then refer back to it rather than re-explaining.
- **Table entries** are shorter and consistent in structure — each says what to do, not why.
- Cut roughly a third of the word count from the previous version.

Archives Act research and guidance note [Not real guidance – internal thinking/proof of concept only]

In-context semantic search, quick hypothesis generation, augmented workflow

Here's the skeleton of a section for best practice guidance for agencies when developing their AI-related records policies, flesh out these sections and polish how they are presented:

Best practice guidance for agencies

1. Ensure staff are approved to use work accounts when using public AI tools and have an appropriate pathway to get funding.
2. Where staff use personal accounts with public tools, ensure guidance on records policy is clearly communicated.
3. Ensure enterprise AI systems have appropriate mechanisms to transfer records into your agency records management system.
4. Ensure records requirements are proportionate
 - a. Your records policy should avoid placing onerous requirements on capturing all AI-related information as records.
 - b. Your NAP policy should allow for disposal of AI-generated information that is not significant – much like how note-taking and written thinking is not needed to be captured.

Show less

Guidance Note – Recordkeeping Obligations When Using AI Tools

Document • Version 5



I kept each item to a single clear point followed by the rationale, matching the style of the staff-facing sections above. The main thing I added beyond your skeleton was the final sentence in item 4 — the practical argument that disproportionate requirements undermine compliance. That felt like the point the whole section is building toward.

Want me to adjust the tone or level of detail on any of the four?

Guidance Note – Recordkeeping When Using AI Tools

Complementary to PSPF Policy Advisory 001-2025 – OFFICIAL Information Use with Generative Artificial Intelligence

Purpose

The PSPF Policy Advisory addresses what information you can put *into* AI tools. This guidance addresses what to do with what comes *out*, to meet your obligations under the Archives Act 1983.

Key points

1. What constitutes a record is determined by legislation and agency policy.

What you need to keep as a record is set out both by legislation and your agency's policies. Records attract obligations under the Archives Act. When deciding whether to keep information as a record, apply the same judgement to AI-generated content as you would to content from any other source, as per applicable legislation and your agency's policies.

2. Records must be managed according to the Archives Act and regulations set by the National Archives of Australia (NAA).

The NAA sets out requirements for management of records. Records can be disposed of under a records authority or your agency's NAP. If you're unsure whether a record can be disposed of, refer to the NAA's records area.

3. AI tools are not recordkeeping systems

If you produce or receive something through AI tools, and your agency's policies require you to keep as a record, it must be kept in an approved records system – the same as any other content. AI tools do not have their own or different recordkeeping obligations. The same rules apply to AI-generated content.

Common scenarios

Scenario	Treatment
Enterprise AI tool (e.g. agency Copilot)	Enterprise AI tools generally will automatically transfer information into your agency's record management system. Even where that information may not need to be kept as a record. As a best practice, avoid deleting information such as chat logs, prompts, and outputs.
Public AI tool on a work account (e.g. approved ChatGPT or Claude account)	Records held in public AI tools must be transferred to your agency's record management system. As chat logs and other information should be transferred regardless of whether they need to be kept as a record. Tools usually allow for easy exporting of all information attached to the account.
Public AI tool on a personal account	Not all chat logs and information connected to a personal account will relate to work matters. As a best practice staff should use work accounts with public tools. Refer to your agency's records policy to determine how to treat information attached to personal accounts.
AI features in everyday tools (e.g. Google AI Overviews)	Transitory material that may include AI-generated outputs are generally not treated differently from any other transitory reference material such as information returned from Google searches or web pages. Follow your agency's record management policies.

If your agency hasn't issued specific guidance on AI and recordkeeping

Raise it with your information management area. The National Archives' *Information management for records created using AI technologies* (January 2025) advises agencies to ensure their NAP policies and records authorities address AI-generated content. Until your agency provides specific direction, apply your existing recordkeeping policies to AI-generated content as you would to any other work content.

Best practice guidance for agencies

Agencies developing or updating policies to address AI and recordkeeping should consider the following:

1. Provide work accounts for public AI tools.

Where your agency permits the use of public AI tools, staff should be able to access them through work accounts rather than relying on personal accounts. This gives the agency visibility over how tools are being used, reduces the risk of records existing only in personal accounts outside agency control, and allows staff to easily adhere to their records management obligations. Agencies should establish a clear pathway for staff to request and receive access, including any necessary funding approvals.

2. Where staff use personal accounts, communicate your records policy clearly.

Some staff will use personal accounts with public AI tools for work purposes, particularly in instances where the agency does not support paying for these tools. Agencies should not assume this is not occurring. Ensure your records policy explicitly addresses this scenario.

3. Ensure enterprise AI systems can transfer records into your records management system.

When procuring or configuring enterprise AI tools, ensure they support the export or transfer of content into your agency's approved records management system. AI tools that do not allow content to be extracted and managed as records create a compliance risk under the Archives Act.

4. Ensure records requirements are proportionate.

Your records policy should not place onerous requirements on staff to capture all AI-related interactions as records. Many interactions with AI tools – such as exploratory prompts, informal thinking, and note-taking – is functionally similar to handwritten notes, rough calculations, or other inconsequential information. This type of low-significance content has never been required to be captured as a record, and AI does not change that.

Your NAP policy should explicitly provide for the routine disposal of AI-generated content that does not meet the threshold for a record. Staff should have confidence that routine use of AI tools will not create an unmanageable recordkeeping burden. Where requirements are disproportionate, staff are less likely to comply with them, or less likely to use the tools at all.

FAS brief

Rote document generation – synthesising from sources



Produce a series of talking points in dot-point form for our FAS based on the content included in this document. Talking points should include, but not be limited to, "if asked" scenarios such as "QUESTION1?", "QUESTION2?", etc, giving answers to these questions based on the content provided. Include some dot points that include a narrative around this situation.



Short story even shorter – It worked but required edits. Still, the foundational information and phrasing sped up the process.

Project timeline

Consolidating needles from haystack, chronologising

[Attached word doc]
Create a chronology with specific dates using the attached document. The document holds email chains. There is a lot of duplicate information, which you may ignore

That's good. Could you please do a more detailed chronology containing only the events between January 1, 2025, and February 4, 2025.

Copilot

Sure, I can help with that. Here is a chronology of the email chains from the document "Email Dump 2.docx":

1. May 13, 2024: [REDACTED]
2. May 22, 2024: [REDACTED]
3. May 24, 2024: finalize her ad
4. May 28, 2024: ago and menti ①.
5. May 28, 2024: [REDACTED]
6. May 29, 2024: general inform
7. May 31, 2024: on the require
8. June 28, 2024: order issue du
9. June 28, 2024: [REDACTED]
10. June 28, 2024: [REDACTED]
11. July 1, 2024: K [REDACTED]
12. July 2, 2024: D documentation
13. July 18, 2024: [REDACTED]
14. July 25, 2024: [REDACTED]
15. September 23, 2024: relationship between
16. September 24, 2024: steps involved ①.
17. January 21, 2025: She [REDACTED]
18. January 31, 2025: She processes before final
19. February 3, 2025: Joh [REDACTED]
20. February 4, 2025: She [REDACTED] fo required ①.

Let me know if you need any further details or assistance!

Copilot

Certainly! Here is a more detailed chronology of the events between January 1, 2025, and February 4, 2025, based on the document "Email Dump 2.docx":

1. January 8, 2025: [REDACTED]
2. January 21, 2025: processes before ①.
3. January 21, 2025: [REDACTED]
4. January 21, 2025: that the draft [REDACTED]
5. January 21, 2025: with [REDACTED]
6. January 24, 2025: [REDACTED] from [REDACTED] longer required ①
7. January 31, 2025: processes before [REDACTED]
8. February 3, 2025: from [REDACTED]
9. February 4, 2025: [REDACTED] required ①.

Let me know if you need any further details or assistance!

Discussion paper [example is 100% deidentified, using dummy text – indicator of length/style/format of output only]

Rote reproduction of facts – human exercises judgement

Here's a bunch of context about the SUBJECT and an upcoming meeting... [copy paste output from Copilot explaining the SUBJECT]... attached is further context and a previous discussion paper. Generate a new discussion paper for the upcoming meeting

Rewrite to be less explicit about SUBJECT, frame it around intent/outcomes rather than specific actions. For each discussion item include high-level context about that item.

Context and purpose

In recent years, extreme weather events have increased in frequency and impact, placing growing pressure on communities, critical infrastructure, and public services. Governments at all levels are responding with targeted recovery measures, but there is a parallel need to strengthen forward-looking resilience—to reduce future costs, improve continuity of essential services, and protect public safety and wellbeing.

The Government has progressed work on an approach to accelerating climate resilience infrastructure across portfolios. This session is intended to consult on and test options to strengthen coordinated investment, uplift delivery capability, and improve the use of data and standards to support resilience outcomes.

This meeting will provide an opportunity for infrastructure operators, state and territory partners, insurers, local government representatives, academics, and community stakeholders to share insights and practical experience. Input from this session will help inform policy design and implementation sequencing.

The discussion will be structured around three themes and sets of open questions. Participants are encouraged to share candid views, examples, and lessons from their own sectors and experience. The themes are intended to frame discussion, not limit it.

Discussion paper [example is 100% deidentified, using dummy text – indicator of length/style/format of output only]

Rote reproduction of facts – human exercises judgement

1. Governance, Accountability & Public Confidence

Climate resilience infrastructure often sits across multiple jurisdictions and portfolios, with complex roles for owners, regulators, planners, and funders. Inconsistent risk tolerances, unclear accountabilities, and limited transparency about prioritisation can slow delivery and reduce confidence that investments are targeted to the highest public-value needs.

International peers are increasingly embedding resilience into infrastructure planning through clearer accountability frameworks, consistent standards, and transparent reporting on risk and adaptation outcomes. Domestically, there is appetite for a more coherent approach that supports better prioritisation, faster decision-making, and clearer public communication on the trade-offs involved.

Illustrative approaches in other jurisdictions (examples for discussion):

- A national resilience standard applied to major infrastructure business cases, with mandatory reporting on climate risk exposure and adaptation measures.
- A cross-government resilience assurance function to review high-impact projects for consistency with resilience objectives and agreed risk thresholds.
- Public dashboards that track resilience investments, expected benefits, and delivery progress in a way that supports accountability and trust.

Questions (for discussion):

- What governance arrangements most effectively support consistent resilience outcomes across portfolios while respecting devolved delivery responsibilities?
- Would a central assurance or coordination mechanism add value (e.g., to review high-impact projects, align standards, or share reusable tools)? What should it not do?
- How can Government improve transparency and public confidence in how resilience investments are prioritised—particularly when trade-offs are unavoidable?
- What are the most practical ways to ensure resilience objectives are embedded early (planning/business case) rather than retrofitted late (delivery/operations)?
- In your experience, what approaches best balance speed of delivery with robust assurance in high-stakes infrastructure decisions?

2. Delivery Workforce & Capability

A step-change in resilience investment and delivery will require capability uplift across the system: climate risk literacy, engineering and project delivery capacity, commercial and procurement expertise, and community engagement capability. Pressure points may include skills shortages in regional areas, constrained professional services capacity, uneven capability across smaller entities, and the operational challenge of delivering adaptation works while maintaining continuity of services.

Resilience also relies on meaningful engagement—particularly with communities disproportionately affected by hazards, and with operators responsible for essential infrastructure (transport, energy, water, telecommunications, health and social infrastructure). Done well, community-informed delivery can improve legitimacy, reduce contestability, and surface practical solutions early.

Illustrative capability levers (examples for discussion):

- A cross-agency resilience delivery community of practice, including shared templates, playbooks, and procurement approaches.
- Targeted capability programs for project sponsors and senior responsible officers on climate risk, adaptation options, and benefits realisation.
- Embedded surge capacity (e.g., a mobile expert pool) to support agencies and regions during planning peaks and rapid post-disaster rebuild windows.

Questions (for discussion):

- What does community-informed resilience delivery look like in practice, particularly for projects with long time horizons and uncertain risk profiles?
- Where are the biggest delivery constraints likely to emerge (skills, approvals, supply chains, governance, funding rules), and what mitigations are most realistic?
- What enabling environment do delivery leaders need (delegations, guidance, assurance, risk appetite, access to experts) to move faster without increasing failure risk?
- Would a cross-government capability and enablement function be useful? If so, what should it focus on (training, surge support, commercial tools, evaluation)?
- What has worked well in your organisation to build and maintain specialist capability in a tight labour market?

Total of 5 discussion segments of this length produced

Easy infographics

ChatGPT images as the new PowerPoint killer - avoid if you need very specific control over slide widgets




AIDE triage model – deciding where and how to act

A flexible decision model to help senior leaders prioritise in a fast-changing AI environment.




Easy infographics

ChatGPT images as the new PowerPoint killer - avoid if you need very specific control over elements



Make an infographic for the following data about the Masked Economist APS podcast viewership/subscribers/guests, leveraging the style and theme and colours from the attached image but also other colours, just make sure it looks good. Data is as of 15 May 2026. Incorporate the below stats:



Modernised Strategic Update Brief

Rethinking how to present information when equipped with AI

- Workflow included Claude, Gemini, ChatGPT, and Copilot
- 2.1 ASL over two days – I had minimal input apart from the infographic.
- None of this involved analysis, summary, or reliance on AI (the humans involved had requisite baseline knowledge)


1. Karthik used Copilot to **semantic search** across files to **compile and categorise information** about the progress our Division has made. Reviewed information to fact check.
2. Karthik used Claude to transfer information into an A4 **HTML dashboard brief** – leveraged our internal branding kit, produced an aesthetic “brief”, non-linear, more like a dashboard.
3. Paul used Claude Code to read the AI CoLab API to **compile attendance data** for 2025-26.
4. I used Gemini/NotebookLM to **semantic search** across our AI CoLab files to find participant quotes about the value of CoLab.
5. I used ChatGPT to **produce an infographic** to go with the brief, merging the data from Claude and the quotes from Gemini



Modernised Strategic Update Brief

Rethinking how to present information when equipped with AI

- Workflow included Claude, Gemini, ChatGPT, and Copilot
 - 2.1 ASL over two days – I had minimal input apart from the infographic.
 - None of this involved analysis, summary, or reliance on AI (the humans involved had requisite baseline knowledge)
1. Karthik used Copilot to **semantic search** across files to **compile and categorise information** about the progress our Division has made. Reviewed information to fact check.
 2. Karthik used Claude to transfer information into an A4 **HTML dashboard brief** – leveraged our internal branding kit, produced an aesthetic “brief”, non-linear, more like a dashboard.
 3. Paul used Claude Code to read the AI CoLab API to **compile attendance data** for 2025-26.
 4. I used Gemini/NotebookLM to **semantic search** across our AI CoLab files to find participant quotes about the value of CoLab.
 5. I used ChatGPT to **produce an infographic** to go with the brief, merging the data from Claude and the quotes from Gemini



AI Strategic Analysis

*The value of huge context inputs to rework base information (not actual analysis), 1600 words in 2 minutes. **2.5 years' worth of Alex Roberts' Weekly Bits of AI Interest blog provided as source material***

Create me a strategic analysis of the implications of these AI developments in the first four months of 2026. Think about a senior leadership audience within the Australian Public Service - so time poor and interested in the 'so what'.

I'll treat this as a leadership "so what" brief rather than a recap: identify the major shifts, implications for APS operating models, and decisions senior leaders may need to bring forward.

The strongest pattern I'm seeing is that "AI adoption" is no longer mainly about individual productivity. By April the centre of gravity had shifted to agents, workflow redesign, cyber exposure, infrastructure dependency, and public legitimacy.

Thought for 37s >

Strategic analysis: AI developments, Jan–Apr 2026

Audience: APS senior leaders | Focus: "so what?"

Note the 'tasking' part of the prompt is small – but the provided context is enormous

Executive bottom line

The first four months of 2026 suggest a shift from **AI as a productivity tool** to **AI as operating infrastructure**. The most strategically significant development is not any single model release. It is the convergence of five things: rapidly improving agentic capability, intensifying cyber risk, rising infrastructure dependency, early labour-market disruption signals, and growing public/political salience.

For the APS, the implication is clear: **AI is moving from optional experimentation to core public administration capability**. The strategic question is no longer "should agencies use AI?" but **how quickly can the APS redesign work, governance, assurance, cyber posture and workforce capability so that AI can be used safely at scale?**

1. Capability acceleration is now operationally relevant

Across the period, frontier capability moved from "impressive demos" toward systems that can work across software, files, workflows and longer tasks. By late April, GPT-5.5 was framed as able to handle messy, multi-part tasks across coding, research, data analysis, documents, spreadsheets and software operation, with less step-by-step management by the user. OpenAI also released workspace agents designed for shared team workflows, while Google released an enterprise agent platform and Deep Research agents positioned as the first step in agentic pipelines. [Weekly bits of AI interest 24 A...](#)

So what for APS leaders: capability is now good enough to challenge the shape of administrative work. This is not just "better drafting" or "faster summarising". It affects how briefs, correspondence, policy research, procurement, grants, compliance, call centres, FOI, records, cyber triage and service navigation could be run.

The danger is that agencies will use 2024 governance and 2010 workflow design for 2026 capability. That produces either underuse, unsafe use, or "islands of productivity" where individuals move faster but enterprise processes do not.

2. The centre of gravity has shifted from chatbots to agents

The most important architectural shift is toward agents that can act, not just answer. The February and March signals include Perplexity Computer, Copilot Tasks, Claude Cowork scheduling, agentic wallets, cloud-hosted managed agents, and public sector examples such as Estonia's view that agentic AI could make policy cycles part of technical infrastructure. [Weekly bits of AI interest 2026...](#)

[Weekly bits of AI interest 2026...](#)

AI Strategic Analysis

*The value of huge context inputs to rework base information (not actual analysis), 1600 words in 2 minutes. **2.5 years' worth of Alex Roberts' Weekly Bits of AI Interest blog provided as source material***

This matters because agents introduce a different risk profile. A chatbot gives advice; an agent may retrieve data, modify files, submit forms, trigger workflows, spend money, contact people, or take actions in connected systems.

So what for APS leaders: "human in the loop" is no longer enough as a slogan. The Singaporean agentic AI governance framework highlighted the need for meaningful checkpoints, unique agent identifiers, clear permissions and ongoing auditing of human oversight. [Weekly bits of AI interest 2026...](#)

The APS needs an **agent governance model**, not just a generative AI use policy. answer: what can agents access, what can they do, when must they stop, who is logged, what is reversible, and what cannot be delegated?

3. Cyber risk has crossed a threshold

Mythos is the clearest discontinuity in the period. The April signal is not simply "AI can chain individually minor weaknesses into system compromise." [Weekly bits of AI interest 20 March 2026.docx](#)

The FreeBSD example in early April reinforced the same point: AI did not merely find developed working exploits, compressing the time, cost and expertise needed for

[Weekly bits of AI interest 3 April 2026.docx](#) The UK AI Safety Institute's evaluation of Mythos also highlighted the importance of cybersecurity basics: patching, access controls, configuration and

[Weekly bits of AI interest 17 April 2026.docx](#)

So what for APS leaders: cyber can no longer be treated as a downstream assurance and cyber uplift are now the same agenda. The APS should assume that adversarial and open models to probe public systems, suppliers and legacy services. The desire to avoid AI, but to use it earlier and more systematically for attack-surface mapping, vulnerability discovery, supplier risk analysis and incident response.

The strategic question is whether the APS can move from **compliance cyber** to **AI at machine speed**.

5. The adoption bottleneck is organisational, not technical

Several signals point in the same direction: AI value is not unlocked by tool access alone. The material on moving from experimentation to transformation noted that pilots often fail to become enterprise systems because risk, governance and workflow have not been redesigned for agents. The UK Open Innovation Team's lessons were especially relevant: adoption is mostly a human problem; the highest-value use cases are workflow-level, not prompt-level; and impact remains because of the jagged frontier. [Weekly bits of AI interest 20 March 2026.docx](#)

So what for APS leaders: the unit of AI transformation is not the individual public servant. It is the workflow. Training people to prompt better is useful, but insufficient. The larger gains will come from redesigning end-to-end work: the handoffs, delegations, records, QA, approvals, escalation and access and assurance processes.

For the APS, likely high-value workflow candidates include correspondence, ministerials, PQs, assessment, procurement analysis, regulatory triage, consultation analysis, call-centre knowledge support, FOI processing, policy synthesis, compliance monitoring and cyber operations.

6. Public trust and workforce legitimacy are now central risks

By April, over 70% of Australians reportedly believed AI would cost jobs, with warnings that this was already hardening into political identity. [Weekly bits of AI interest 10 April 2026.docx](#) The same period included reports of AI exposure across occupations, concerns about early-career roles, staff reductions in major firms, and research suggesting uneven impacts across the labour market. [Weekly bits of AI interest 20 March 2026.docx](#)

Within the APS, the Assistant Minister's March message is strategically important: AI is not only theoretically, every public servant must learn to use it, and "AI wrote it, not me" is not an excuse. It remains the responsibility of the public servant. [Weekly bits of AI interest 20 March 2026.docx](#)

So what for APS leaders: AI adoption must be seen as a way to make government more accountable, not just a means for covert headcount reduction. Staff will not engage if AI is seen as a way to make government less accountable. [Weekly bits of AI interest 20 March 2026.docx](#)

The APS needs a credible compact: public servants will be expected to use AI, but agencies will invest in capability, redesign work responsibly, protect accountability, and be transparent about where AI is used.

7. Demand on government may rise, not fall

A recurring signal is that AI reduces friction. That is good for citizens, but it can expose weaknesses in policy and delivery systems that depend on friction, ambiguity or complexity to manage demand. AI-assisted applicants will be better at finding entitlements, writing submissions, lodging claims, appealing decisions and navigating processes. The February material explicitly warned that AI-powered public service can produce more service requests, not fewer. [Weekly bits of AI interest 2026...](#)

The February signal on Australian patent filings is a concrete local example: self-filed provisional applications nearly tripled, with AI potentially bringing new and problematic users into complex systems. [Weekly bits of AI interest 2026...](#)

So what for APS leaders: AI adoption cannot only focus on internal efficiency. Agencies also need to prepare for **AI-mediated citizens**. That means clearer rules, better forms, stronger fraud controls, scalable triage, accessible explanations, and policy settings that do not rely on confusion as a rationing mechanism.

Strategic decision agenda for APS leaders

1. **Treat AI as core operating capability.** Move from experimentation portfolios to priority workflows with named executive owners, delivery pathways and measurable service outcomes.
2. **Build agent-ready governance.** Establish clear rules for permissions, identity, logs, records, approvals, escalation, reversibility and prohibited actions for AI agents.
3. **Accelerate cyber uplift using AI.** Prioritise attack-surface mapping, supplier visibility, code review, patching, logging and AI-assisted vulnerability detection.
4. **Redesign work, not just training.** Pair AI literacy with workflow redesign, QA redesign and role redesign. The question is not "can staff use AI?" but "has the work been made AI-compatible?"
5. **Develop a practical sovereignty posture.** Decide which workloads require domestic execution, which require frontier access, which require open-source options, and where procurement must preserve portability.
6. **Prepare for labour-market and legitimacy impacts.** Be explicit that AI adoption is about better services and better work, while building credible transition pathways for exposed roles and early-career capability.
7. **Plan for demand amplification.** Review high-volume public-facing systems for what happens when citizens, firms and representatives use AI agents to interact with government at scale.

The "so what" in one sentence

The APS does not need to chase every model release, but it does need to move faster on the institutional plumbing: accountable agents, AI-ready workflows, cyber resilience, workforce capability, sovereign options and public trust. [Weekly bits of AI interest 20 March 2026.docx](#)

Pitfalls

Beware of these common errors – LLMs are not all-knowing and all-powerful



Magic prompts

Your prompt isn't a magic phrase that retrieves your NPP from the ether. Your NPP doesn't exist – you build it iteratively with the LLM. Stop asking for prompt libraries!



Implicit context

Not all words are said, not all thoughts are written, and not all context can be imparted. Some tasks just have way too much implicit context that the LLM doesn't have – particularly with spatial reasoning.



Hammer → nail

Don't fall for the hammer nail problem. LLMs are a tool in the toolbox, not the answer to everything. I use LLMs probably less than you imagine I do.



Infant capabilities

LLMs are pattern machines and still in their relative infancy – we still have logical and contextual advantage. Never defer, always supplement thinking.

Perspectives on LLMs




Augmentation

You're missing half the gain by only offloading mundane tasks. Use it to achieve more and uplift the quality of work – **this way, you retain brain cells and employability**, and it's just more fun.



Voluntary uplift

The same tools that allowed us to offload can allow us to uplift. Just as **machines offloaded our physical activity, they can uplift it too**. Education is critical so that we can all use it to uplift.



Adaptation

Most of us can't farm or track seasonal changes anymore – but innovation allows us to move on and unlock new tasks/thinking, while providing us a way back (just Google Search how to farm).



Chatbots are limited

LLMs right now are trained as chatbots – digital friends that are **agreeable and conversational**. As demand for rigour goes up, this may change. Other tools exist (NotebookLM for example is more grounded).

Creativity

- With these tools, there's a lot of **creativity** on your part to find ways to use them. It also takes **time and effort to experiment** – and for many of us we're time-poor. **It doesn't come automatically!**
- Don't withhold learning these tools – they keep getting better with time, but the principles of using them remain the same.





Open to the floor

If you're APS – Sign up to GovAI Chat Beta trial.

Download ChatGPT/Claude/Gemini to your phone, it's free! Learn by doing – it is a tool like any other

Visit events.aicolab.org if you want to learn more and get involved with other AI CoLab events!

